

Mean, median & mode imputations

DEALING WITH MISSING DATA IN PYTHON



Suraj Donthi

Deep Learning & Computer Vision
Consultant

Basic imputation techniques

- constant (e.g. 0)
- mean
- median
- mode or most frequent

Mean Imputation

```
from sklearn.impute import SimpleImputer  
diabetes_mean = diabetes.copy(deep=True)  
mean_imputer = SimpleImputer(strategy='mean')
```

Mean Imputation

```
from sklearn.impute import SimpleImputer
diabetes_mean = diabetes.copy(deep=True)
mean_imputer = SimpleImputer(strategy='mean')
diabetes_mean.iloc[:, :] = mean_imputer.fit_transform(diabetes_mean)
```

Median imputation

```
diabetes_median = diabetes.copy(deep=True)
median_imputer = SimpleImputer(strategy='median')
diabetes_median.iloc[:, :] = median_imputer.fit_transform(diabetes_median)
```

Mode imputation

```
diabetes_mode = diabetes.copy(deep=True)
mode_imputer = SimpleImputer(strategy='most_frequent')
diabetes_mode.iloc[:, :] = mode_imputer.fit_transform(diabetes_mode)
```

Imputing a constant

```
diabetes_constant = diabetes.copy(deep=True)
constant_imputer = SimpleImputer(strategy='constant', fill_value=0)
diabetes_constant.iloc[:, :] = constant_imputer.fit_transform(diabetes_constant)
```

Scatterplot of imputation

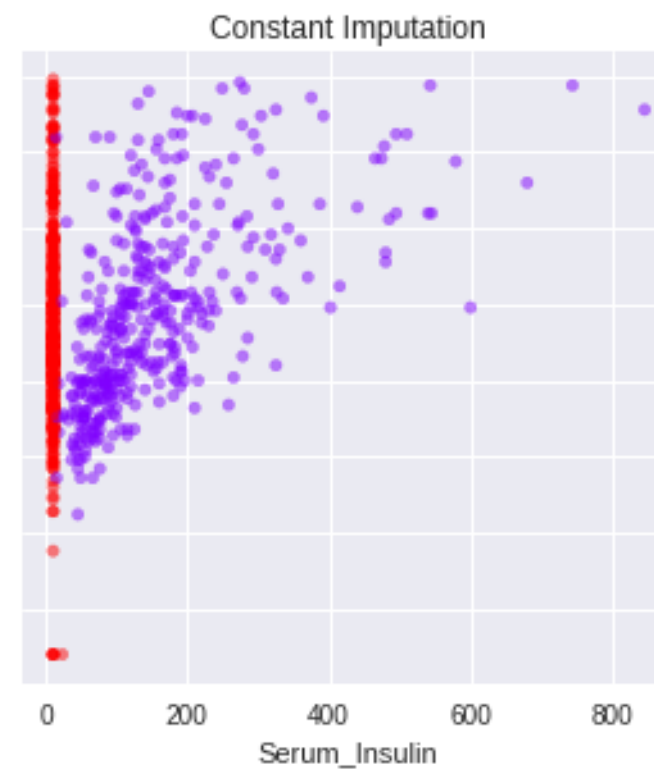
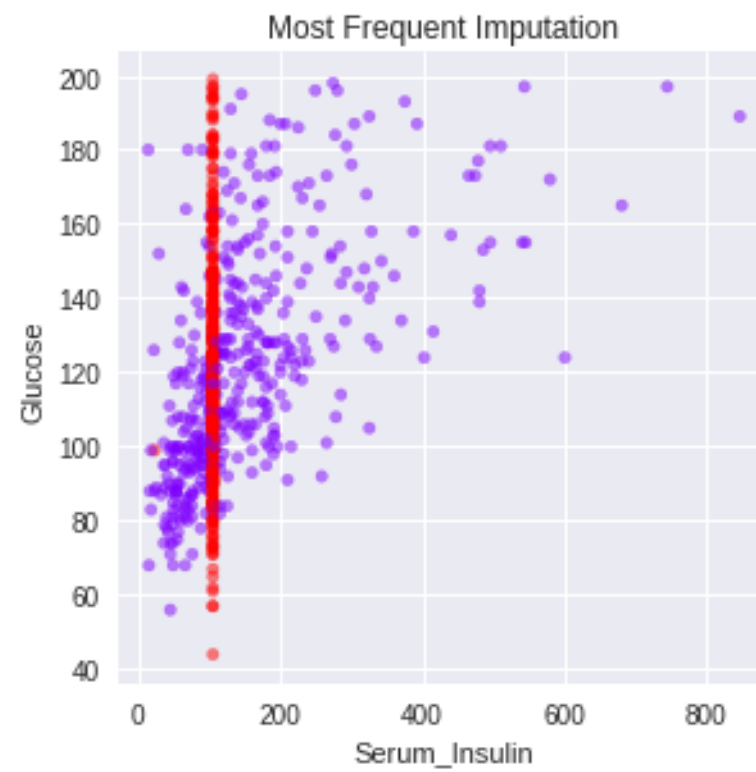
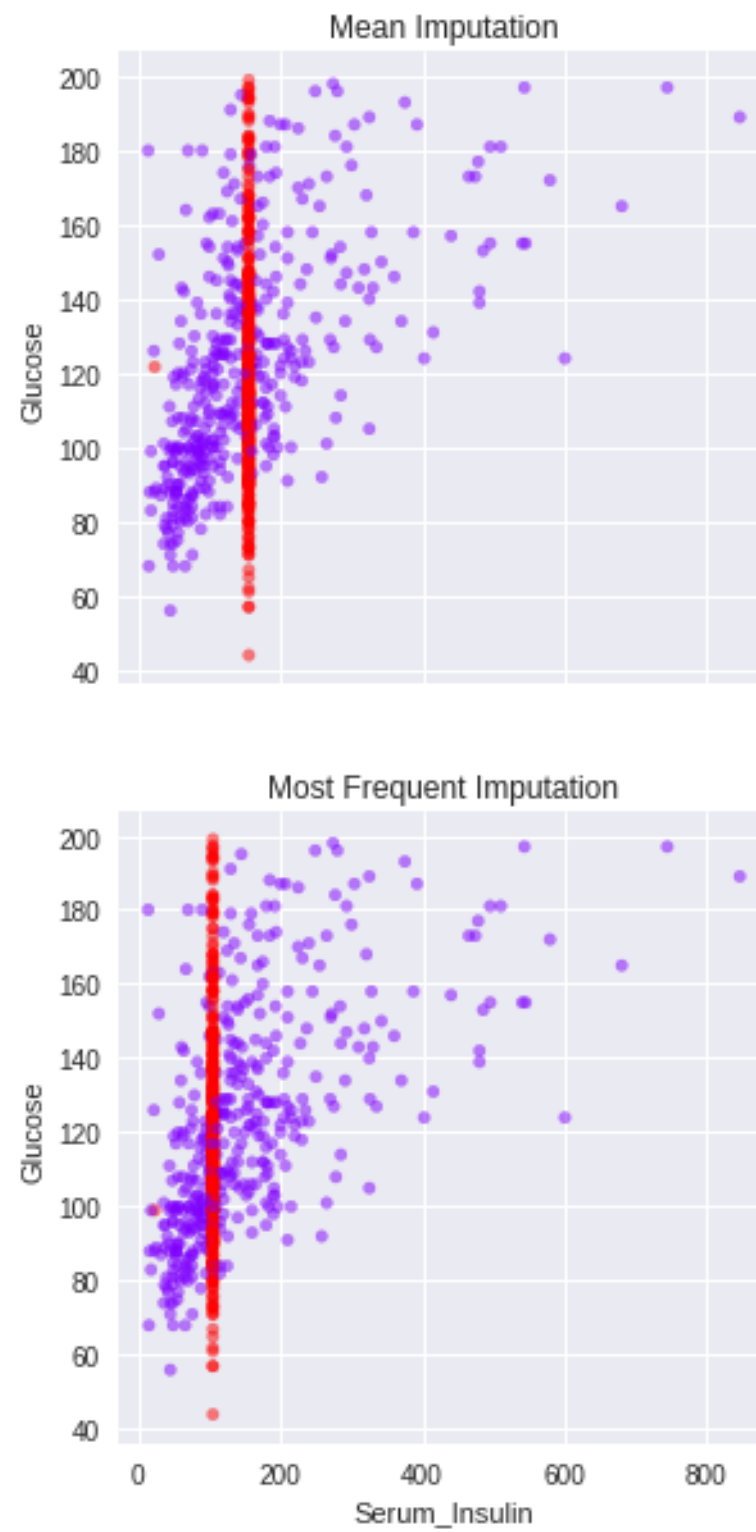
```
nullity = diabetes['Serum_Insulin'].isnull()+diabetes['Glucose'].isnull()
```

```
diabetes_mean.plot(x='Serum_Insulin', y='Glucose', kind='scatter', alpha=0.5,  
                  c=nullity, cmap='rainbow', title='Mean Imputation')
```



Visualizing imputations

```
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(10, 10))
nullity = diabetes['Serum_Insulin'].isnull()+diabetes['Glucose'].isnull()
imputations = {'Mean Imputation': diabetes_mean,
               'Median Imputation': diabetes_median,
               'Most Frequent Imputation': diabetes_mode,
               'Constant Imputation': diabetes_constant}
for ax, df_key in zip(axes.flatten(), imputations):
    imputations[df_key].plot(x='Serum_Insulin', y='Glucose', kind='scatter',
                             alpha=0.5, c=nullity, cmap='rainbow', ax=ax,
                             colorbar=False, title=df_key)
```



Summary

You learned to

- Impute with statistical parameters like mean, median and mode
- Graphically compare the imputations
- Analyze the imputations

Let's practice!

DEALING WITH MISSING DATA IN PYTHON

Imputing time-series data

DEALING WITH MISSING DATA IN PYTHON



Suraj Donthi

Deep Learning & Computer Vision
Consultant

Airquality Dataset

```
import pandas as pd
airquality = pd.read_csv('air-quality.csv', parse_dates='Date',
                        index_col='Date')

airquality.head()
```

	Ozone	Solar	Wind	Temp
Date				
1976-05-01	41.0	190.0	7.4	67
1976-05-02	36.0	118.0	8.0	72
1976-05-03	12.0	149.0	12.6	74
1976-05-04	18.0	313.0	11.5	62
1976-05-05	NaN	NaN	14.3	56

Airquality Dataset

```
airquality.isnull().sum()
```

```
Ozone      37  
Solar       7  
Wind        0  
Temp        0  
dtype: int64
```

```
airquality.isnull.mean() * 100
```

```
Ozone      24.183007  
Solar       4.575163  
Wind        0.000000  
Temp        0.000000  
dtype: float64
```

The `.fillna()` method

The attribute `method` in `.fillna()` can be set to

- `'ffill'` or `'pad'`
- `'bfill'` or `'backwardfill'`

Ffill method

- Replace NaN s with last observed value
- pad is the same as 'ffill'

```
airquality.fillna(method='ffill', inplace=True)
```

```
airquality['Ozone'][30:40]
```

Date	Ozone
1976-05-31	37.0
1976-06-01	NaN
1976-06-02	NaN
1976-06-03	NaN
1976-06-04	NaN
1976-06-05	NaN
1976-06-06	NaN
1976-06-07	29.0
1976-06-08	NaN
1976-06-09	71.0

```
airquality.fillna(method='ffill',  
                  inplace=True)
```

```
airquality['Ozone'][30:40]
```

Date	Ozone
1976-05-31	37.0
1976-06-01	37.0
1976-06-02	37.0
1976-06-03	37.0
1976-06-04	37.0
1976-06-05	37.0
1976-06-06	37.0
1976-06-07	29.0
1976-06-08	29.0
1976-06-09	71.0

Bfill method

- Replace NaN s with next observed value
- `backfill` is the same as `'bfill'`

```
df.fillna(method='bfill', inplace=True)
```

```
airquality['Ozone'][30:40]
```

Date	Ozone
1976-05-31	37.0
1976-06-01	NaN
1976-06-02	NaN
1976-06-03	NaN
1976-06-04	NaN
1976-06-05	NaN
1976-06-06	NaN
1976-06-07	29.0
1976-06-08	NaN
1976-06-09	71.0

```
airquality.fillna(method='bfill',  
                  inplace=True)
```

```
airquality['Ozone'][30:40]
```

Date	Ozone
1976-05-31	37.0
1976-06-01	29.0
1976-06-02	29.0
1976-06-03	29.0
1976-06-04	29.0
1976-06-05	29.0
1976-06-06	29.0
1976-06-07	29.0
1976-06-08	71.0
1976-06-09	71.0

The `.interpolate()` method

- The `.interpolate()` method extends the sequence of values to the missing values

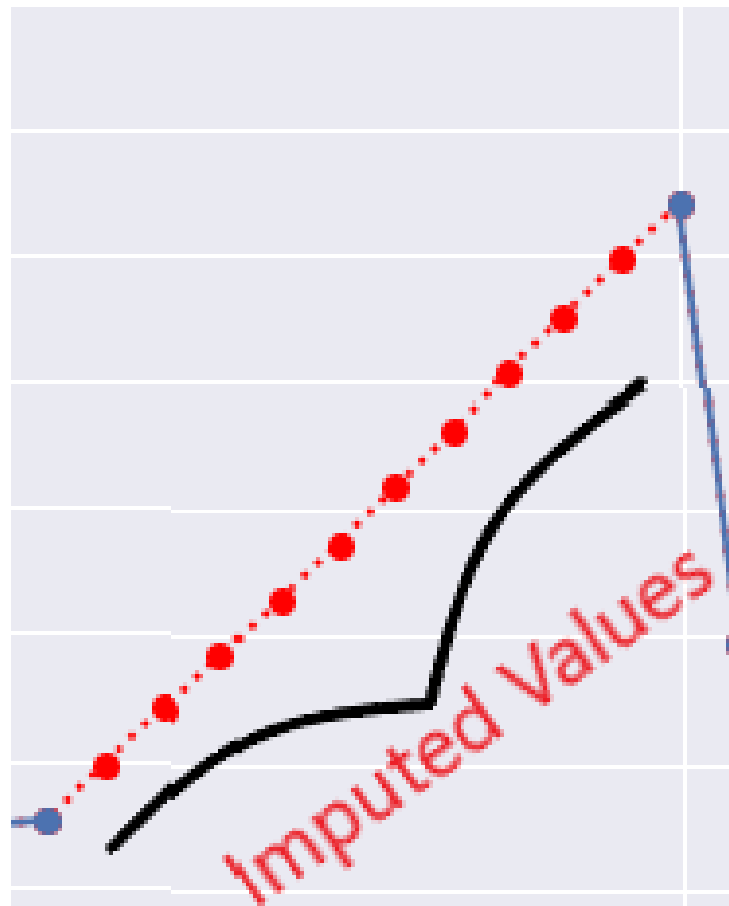
The attribute `method` in `.interpolate()` can be set to

- `'linear'`
- `'quadratic'`
- `'nearest'`

Linear interpolation

- Impute linearly or with equidistant values

```
df.interpolate(method='linear', inplace=True)
```



```
airquality['Ozone'][30:40]
```

Date	Ozone
1976-05-31	37.0
1976-06-01	NaN
1976-06-02	NaN
1976-06-03	NaN
1976-06-04	NaN
1976-06-05	NaN
1976-06-06	NaN
1976-06-07	29.0
1976-06-08	NaN
1976-06-09	71.0

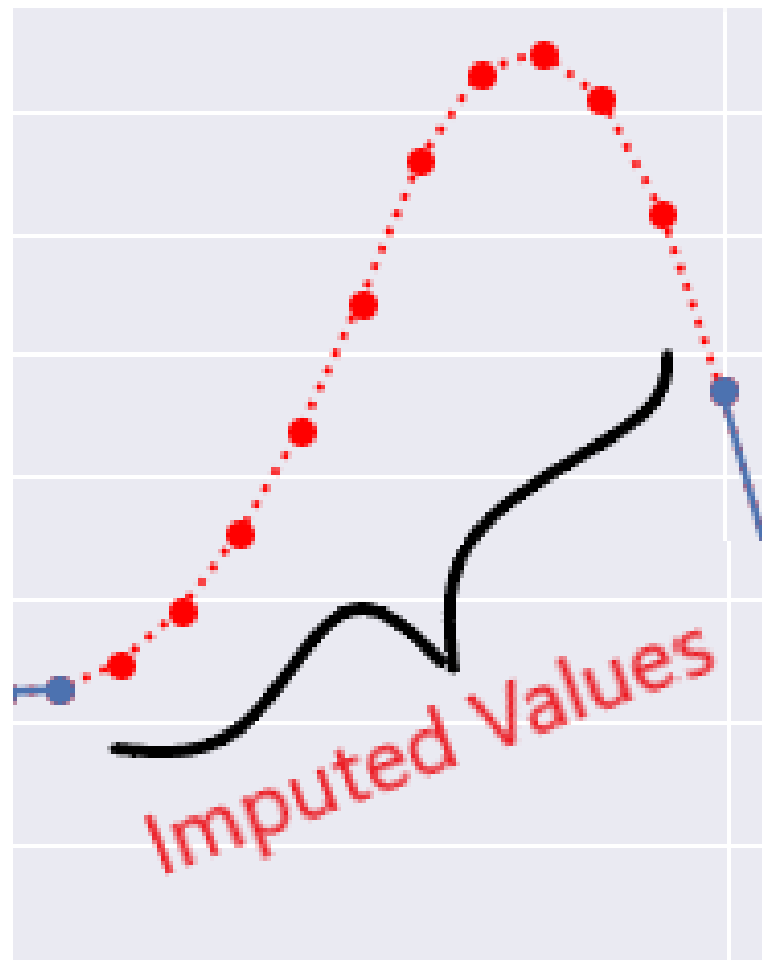
```
airquality.interpolate(  
    method='linear', inplace=True)  
airquality['Ozone'][30:40]
```

Date	Ozone
1976-05-31	37.0
1976-06-01	35.9
1976-06-02	34.7
1976-06-03	33.6
1976-06-04	32.4
1976-06-05	31.3
1976-06-06	30.1
1976-06-07	29.0
1976-06-08	50.0
1976-06-09	71.0

Quadratic interpolation

- Impute the values quadratically

```
df.interpolate(method='quadratic', inplace=True)
```




```
airquality['Ozone'][30:39]
```

Date	Ozone
1976-05-31	37.0
1976-06-01	NaN
1976-06-02	NaN
1976-06-03	NaN
1976-06-04	NaN
1976-06-05	NaN
1976-06-06	NaN
1976-06-07	29.0
1976-06-08	NaN

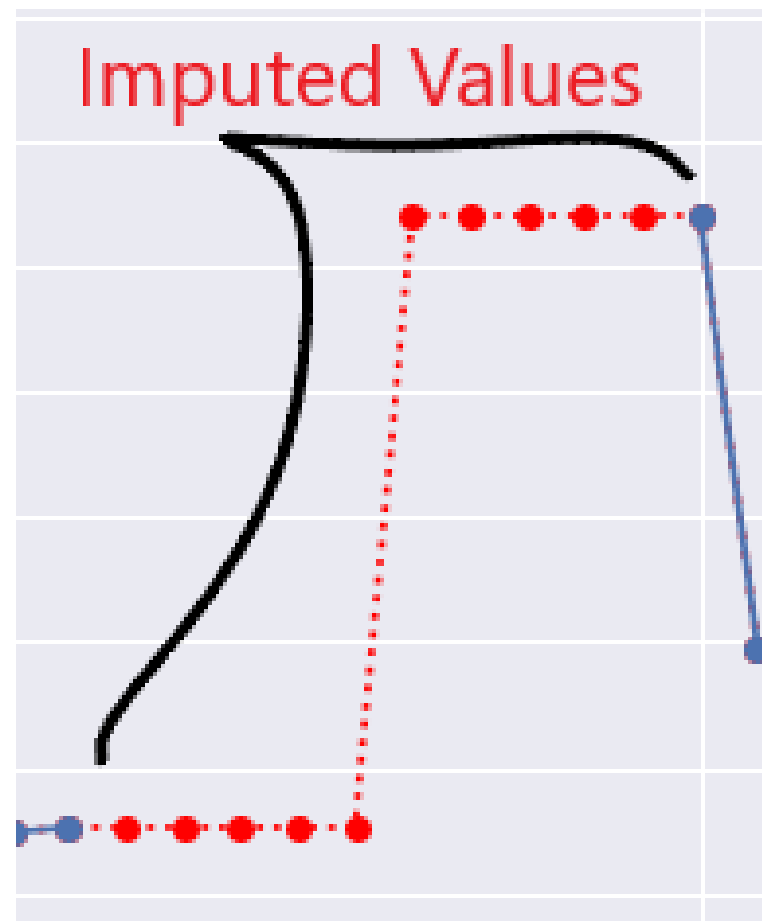
```
airquality.interpolate(  
    method='quadratic', inplace=True)  
airquality['Ozone'][30:39]
```

Date	Ozone
1976-05-31	37.0
1976-06-01	-38.4
1976-06-02	-79.4
1976-06-03	-85.9
1976-06-04	-62.4
1976-06-06	-2.8
1976-06-07	29.0
1976-06-08	62.2

Nearest value imputation

- Impute with the nearest observable value

```
df.interpolate(method='nearest', inplace=True)
```



```
airquality['Ozone'][30:39]
```

Date	Ozone
1976-05-31	37.0
1976-06-01	NaN
1976-06-02	NaN
1976-06-03	NaN
1976-06-04	NaN
1976-06-05	NaN
1976-06-06	NaN
1976-06-07	29.0
1976-06-08	NaN

```
airquality.interpolate(  
    method='nearest', inplace=True)  
airquality['Ozone'][30:39]
```

Date	Ozone
1976-05-31	37.0
1976-06-01	37.0
1976-06-02	37.0
1976-06-03	37.0
1976-06-04	29.0
1976-06-05	29.0
1976-06-06	29.0
1976-06-07	29.0
1976-06-08	29.0

Let's practice!

DEALING WITH MISSING DATA IN PYTHON

Visualizing time-series imputations

DEALING WITH MISSING DATA IN PYTHON

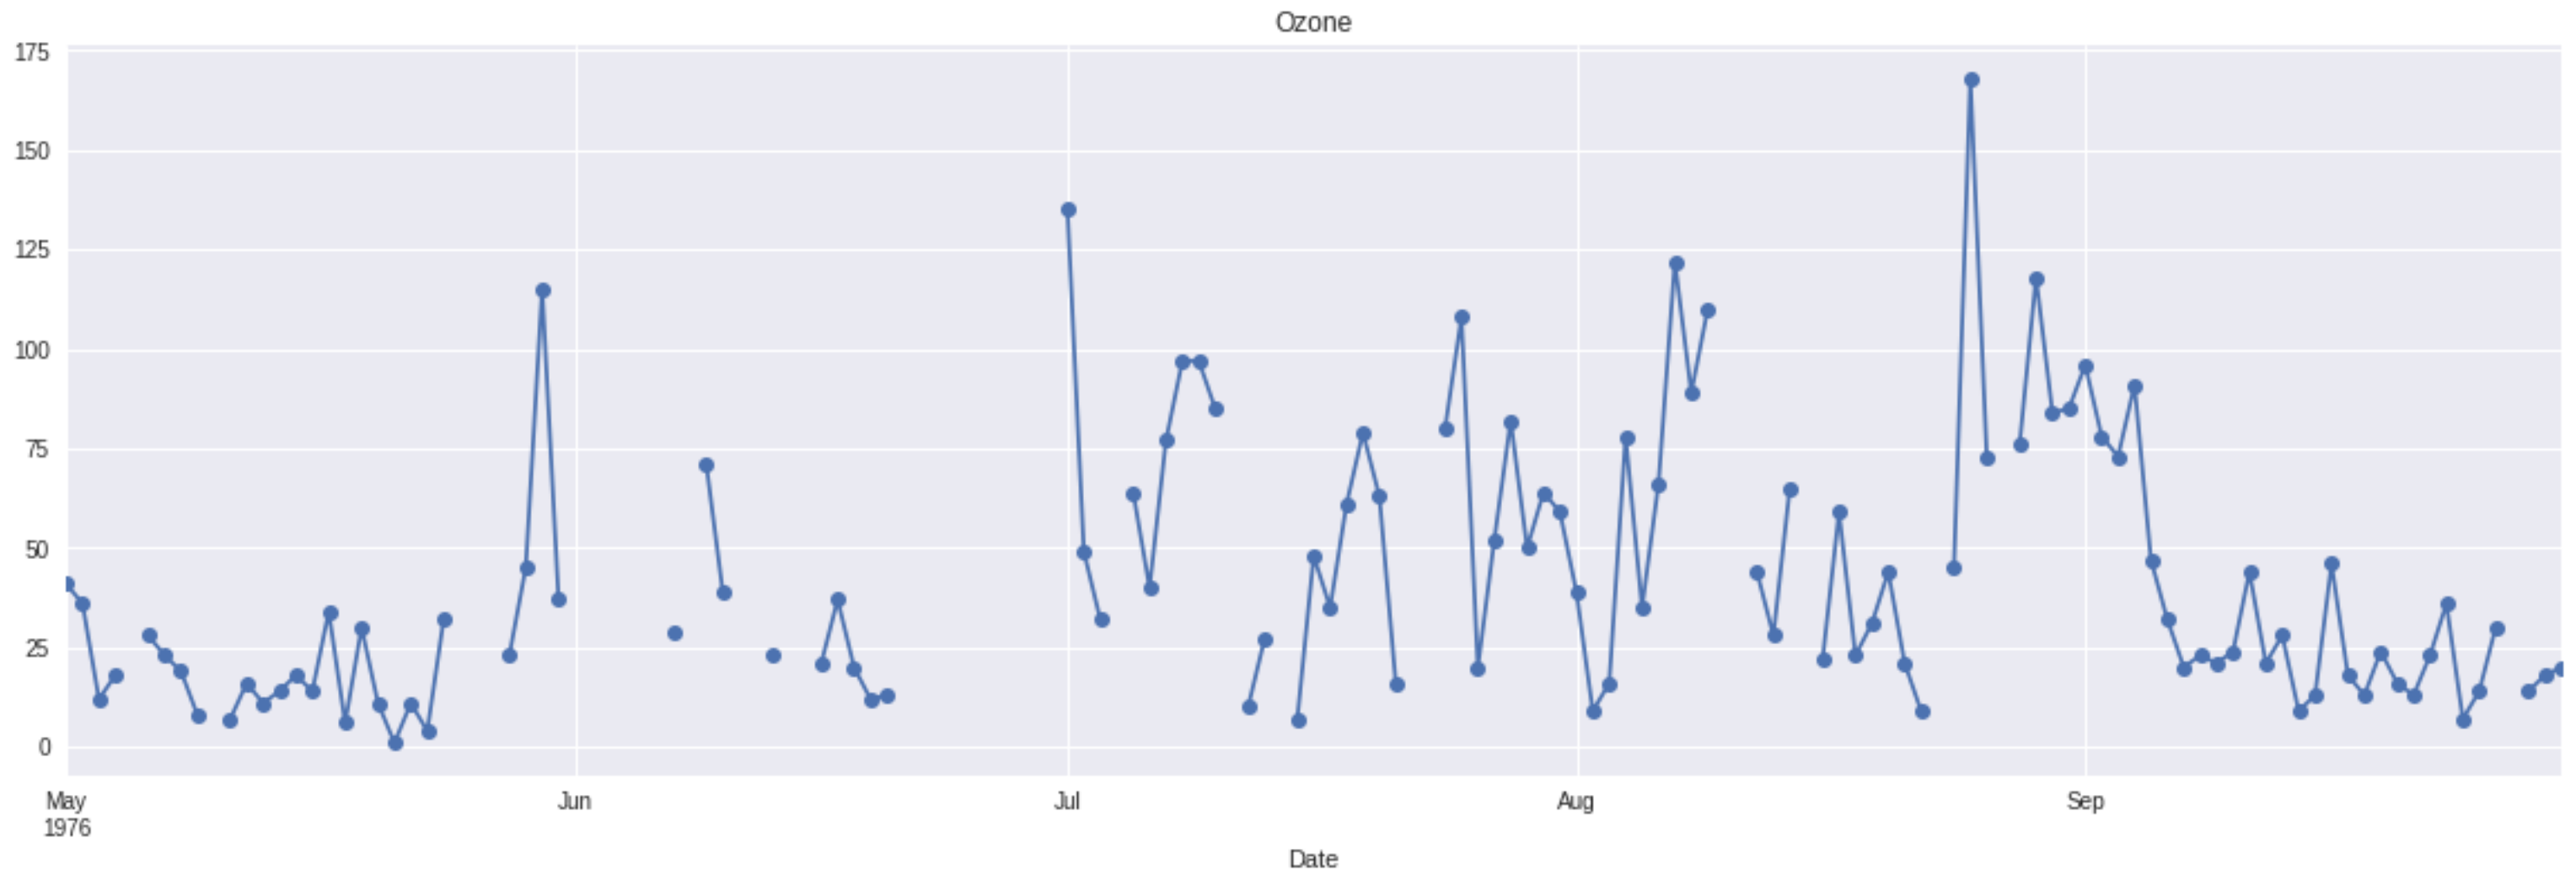


Suraj Donthi

Deep Learning & Computer Vision
Learning

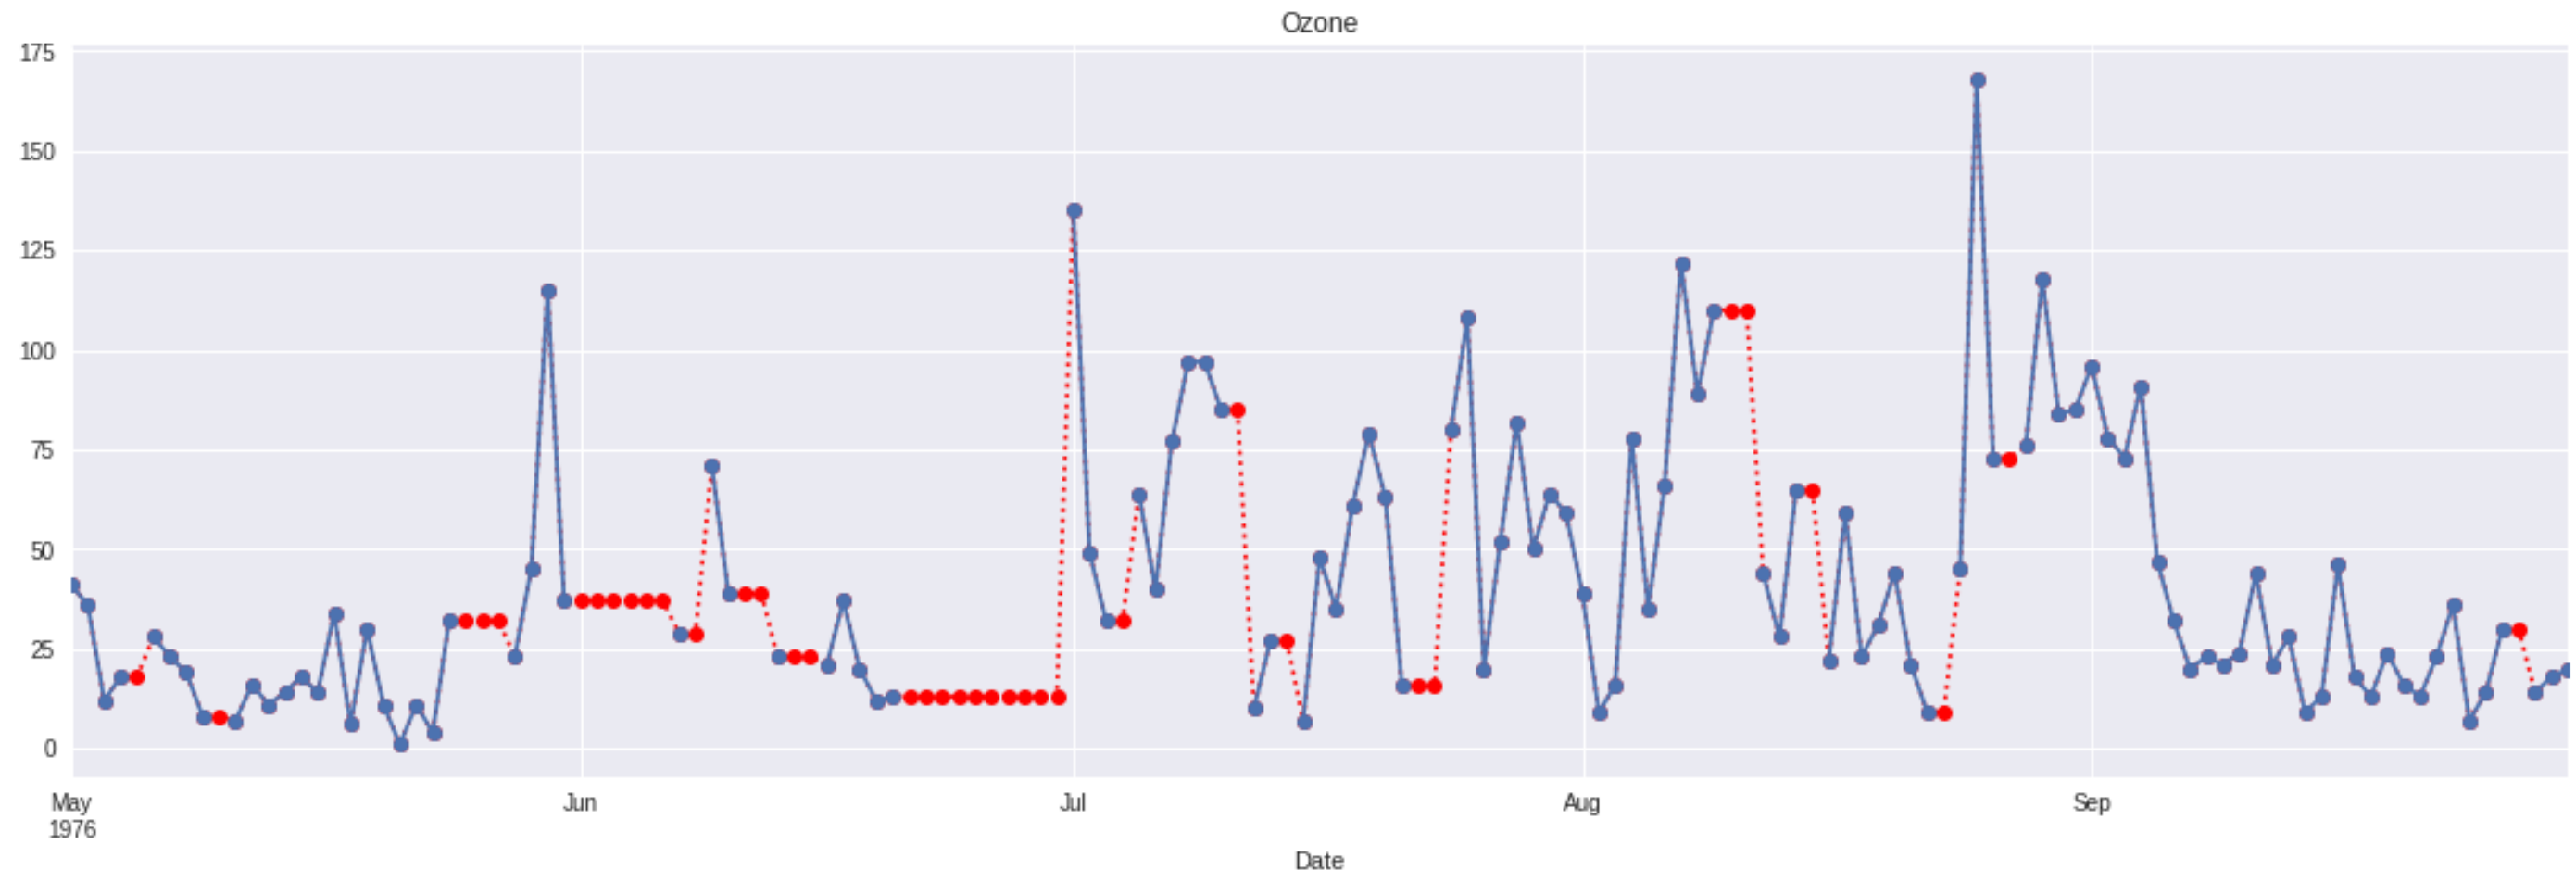
Air quality time-series plot

```
airquality['Ozone'].plot(title='Ozone', marker='o', figsize=(30, 5))
```



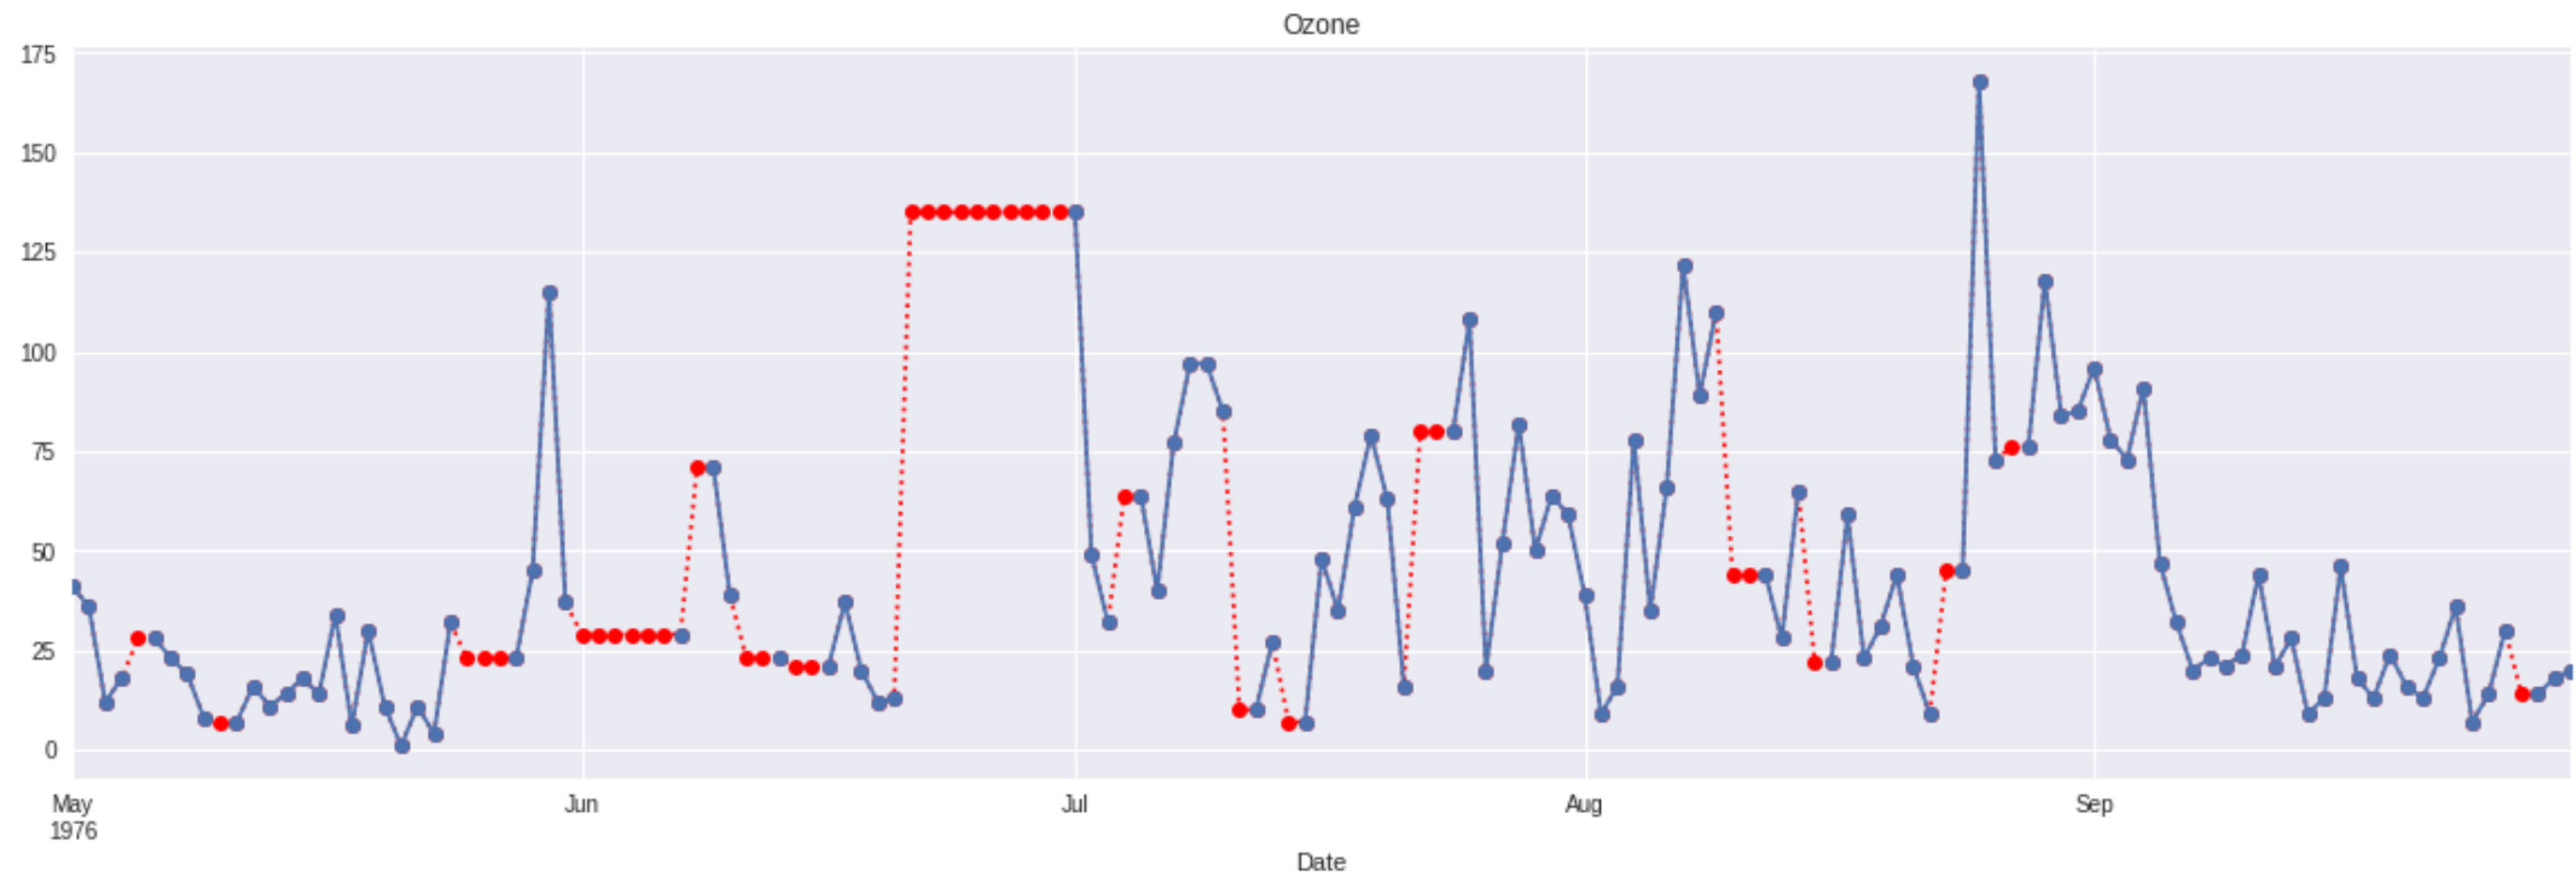
Ffill Imputation

```
ffill_imp['Ozone'].plot(color='red', marker='o', linestyle='dotted', figsize=(30, 5))  
airquality['Ozone'].plot(title='Ozone', marker='o')
```



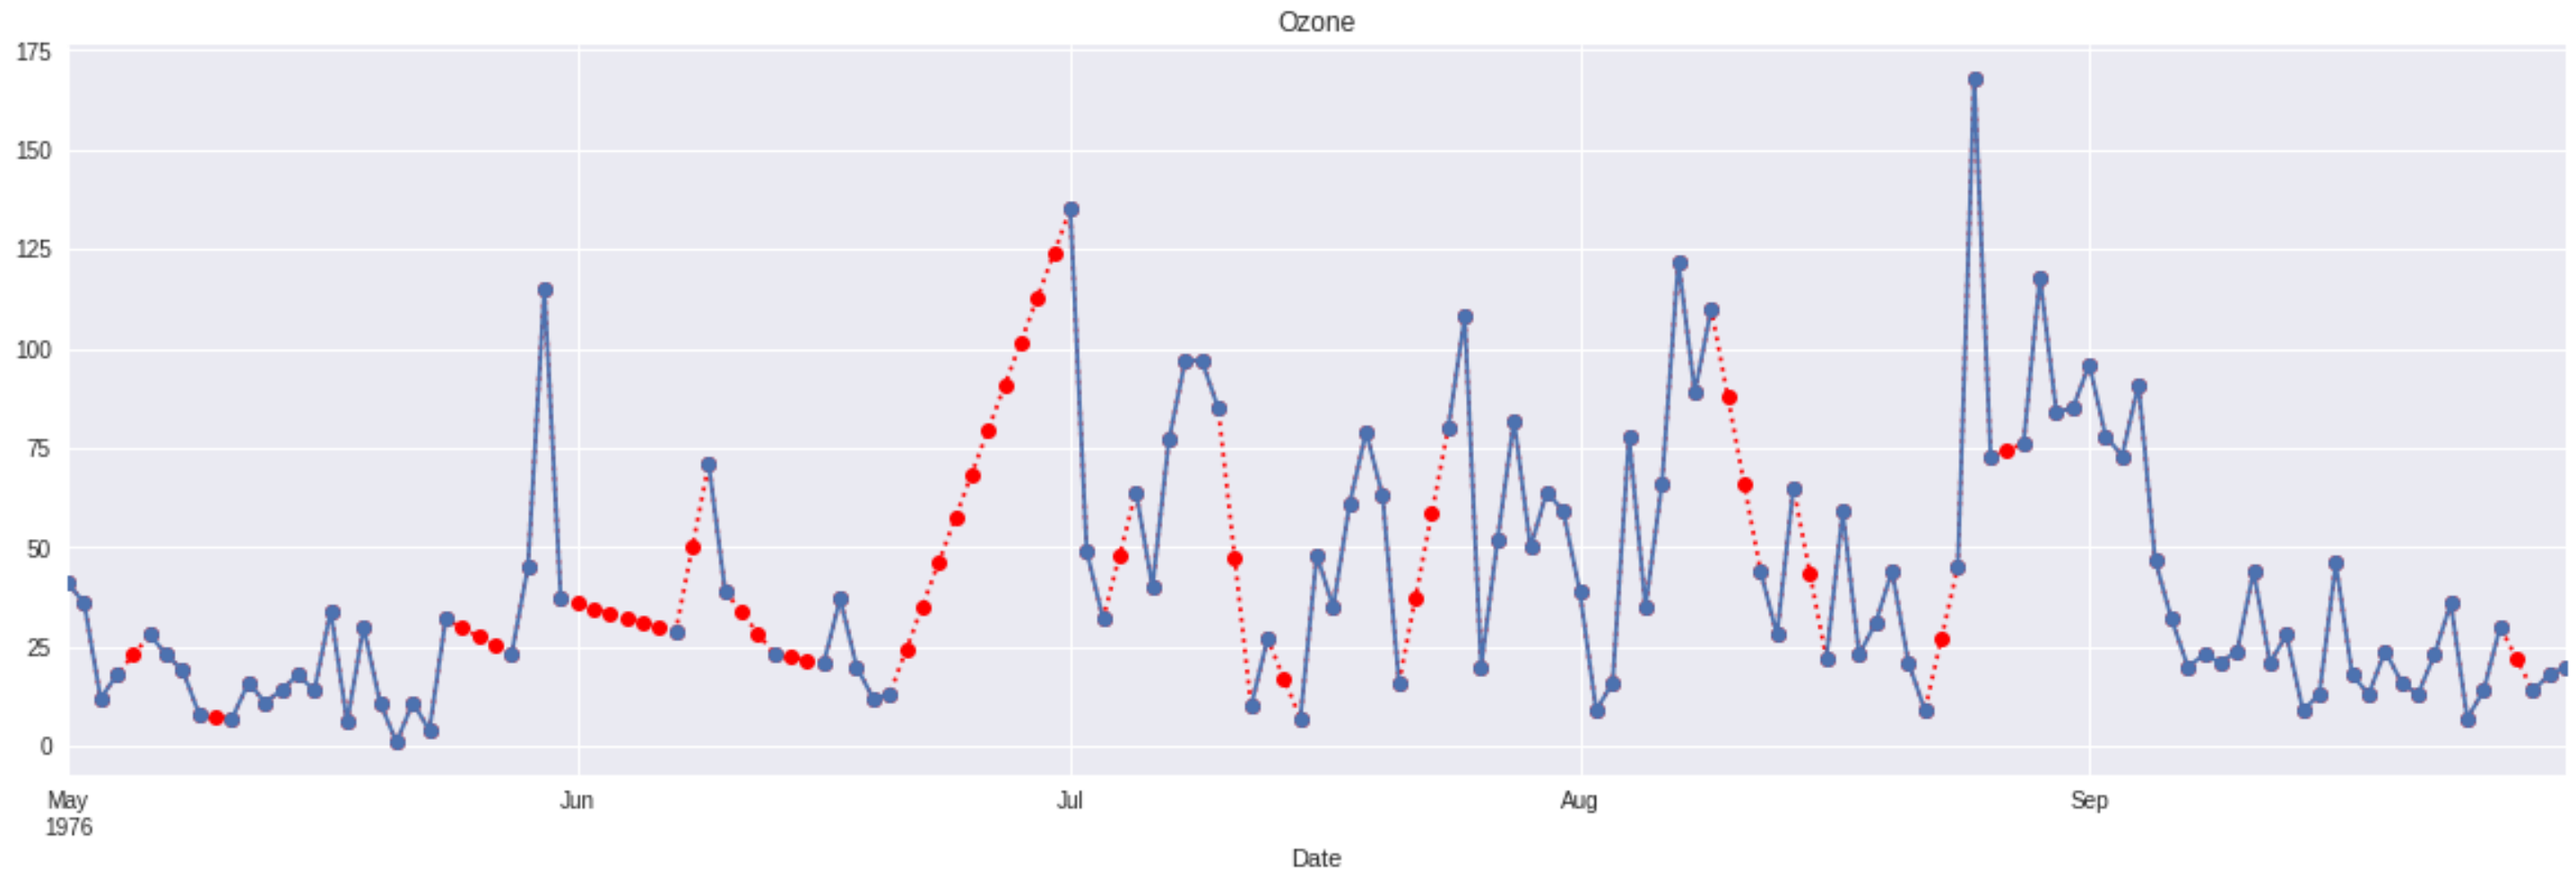
Bfill Imputation

```
bfill_imp['Ozone'].plot(color='red', marker='o', linestyle='dotted', figsize=(30, 5))  
airquality['Ozone'].plot(title='Ozone', marker='o')
```



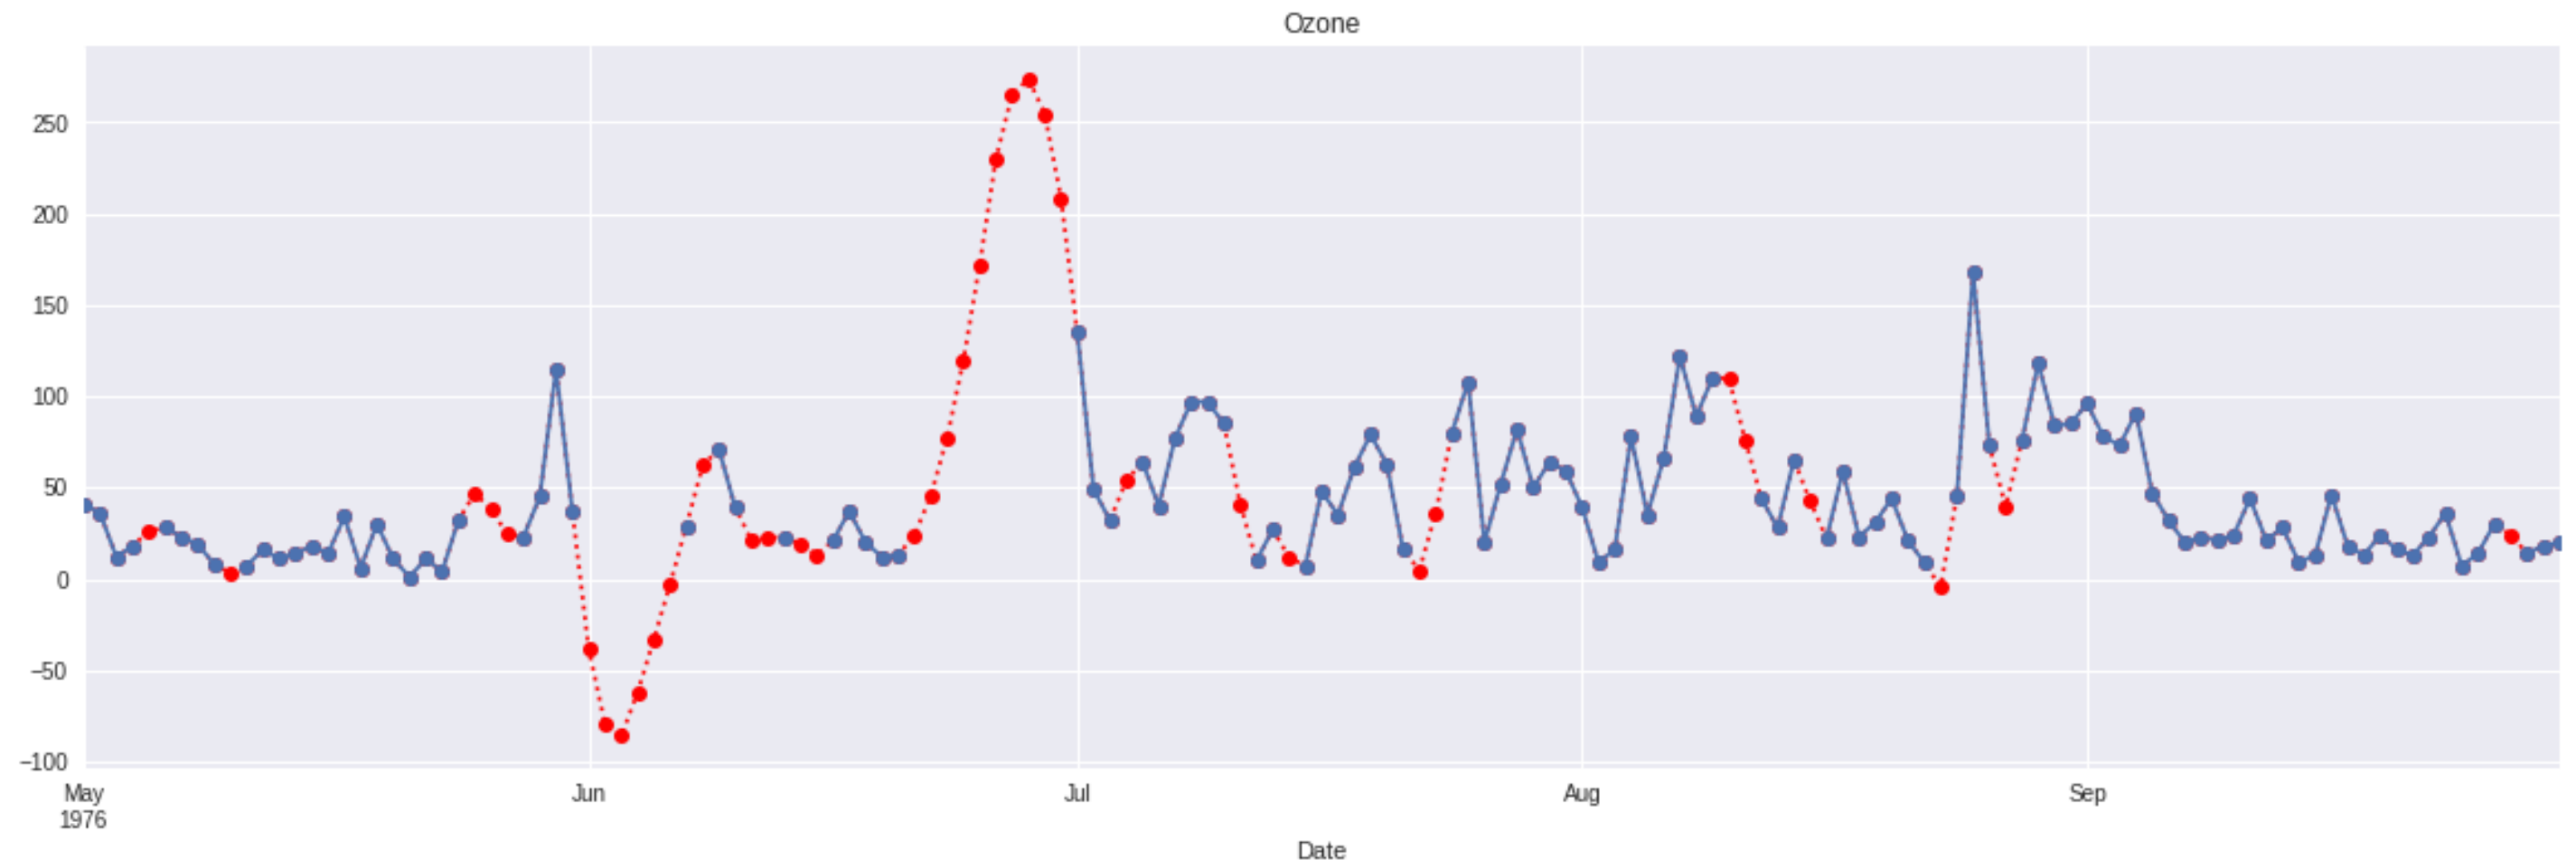
Linear Interpolation

```
linear_interp['Ozone'].plot(color='red', marker='o', linestyle='dotted', figsize=(30, 5))  
airquality['Ozone'].plot(title='Ozone', marker='o')
```



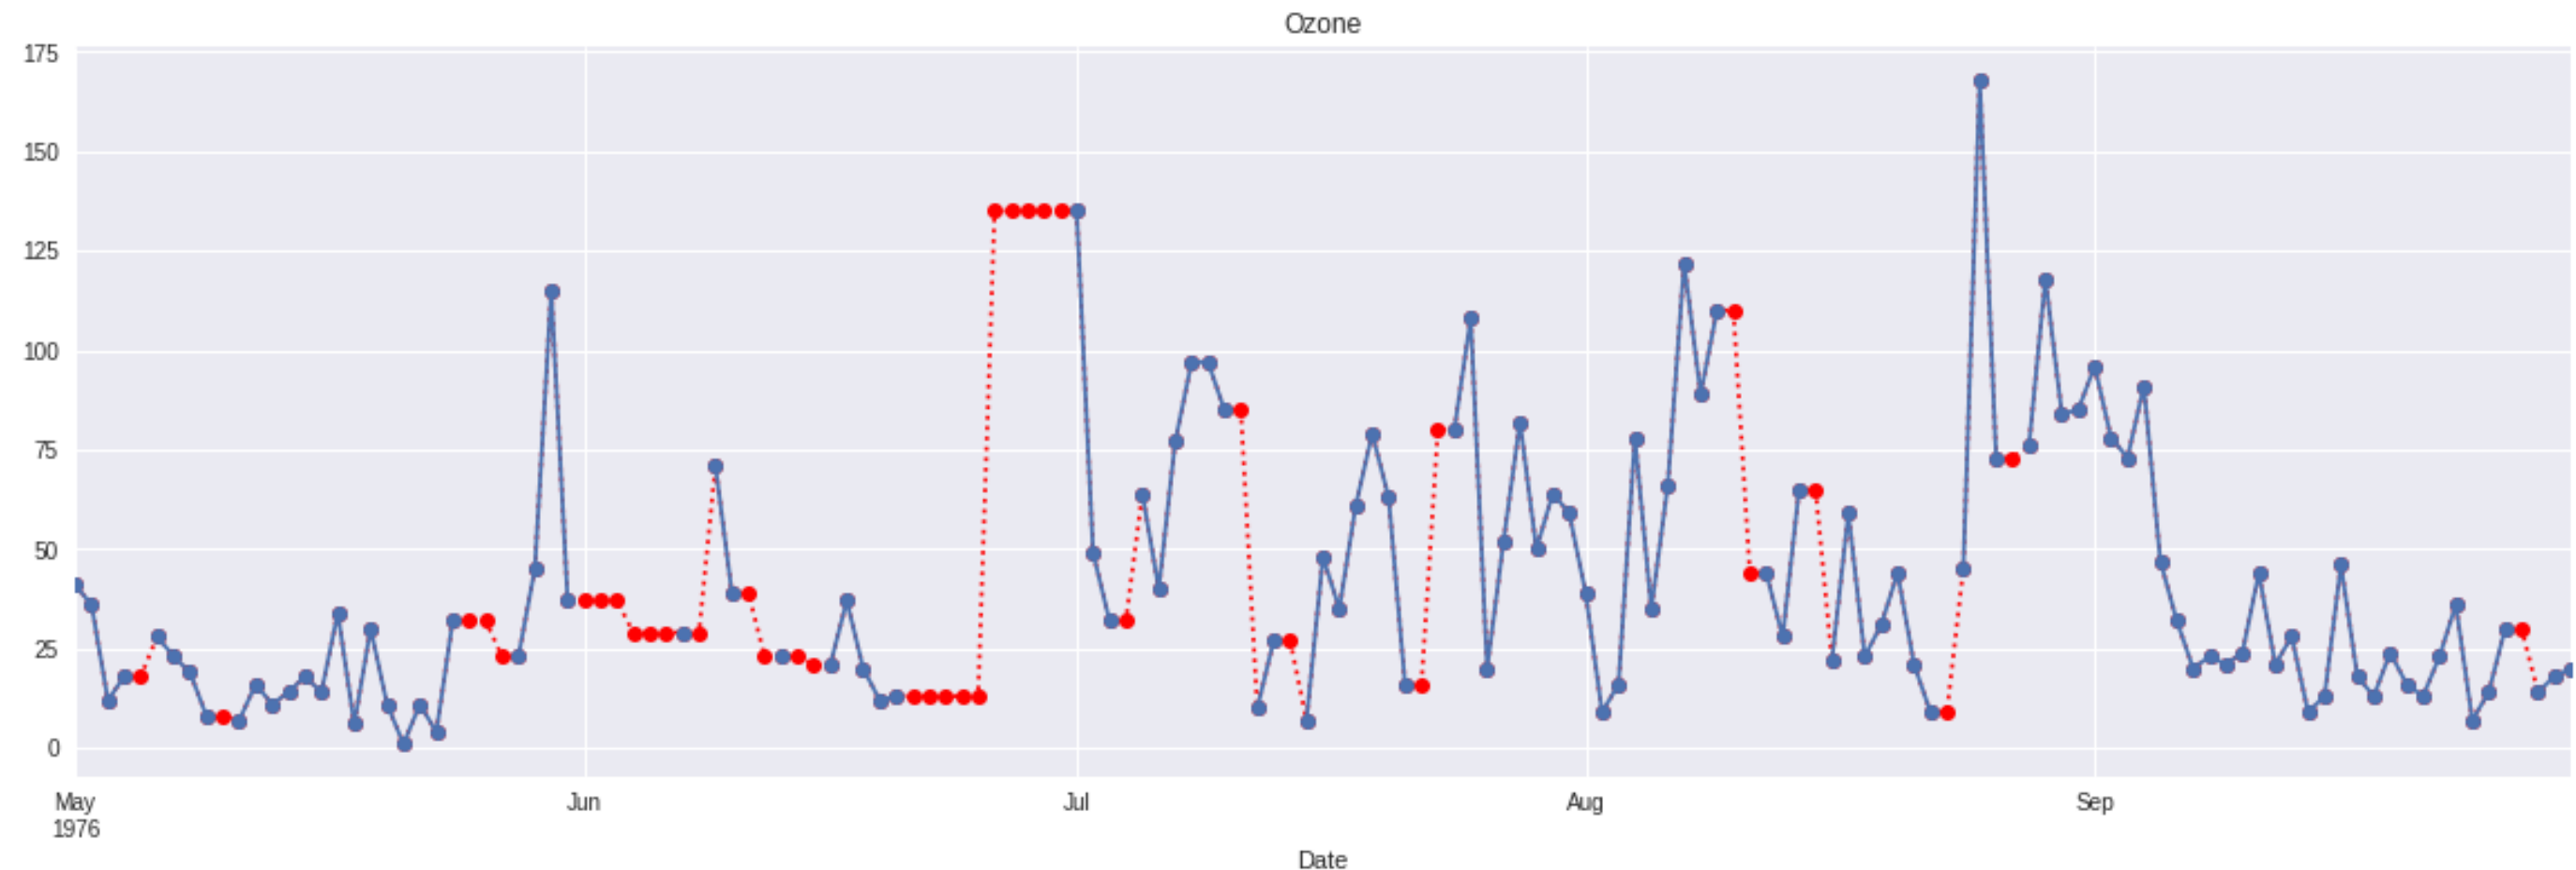
Quadratic Interpolation

```
quadratic_interp['Ozone'].plot(color='red', marker='o', linestyle='dotted', figsize=(30, 5))  
airquality['Ozone'].plot(title='Ozone', marker='o')
```



Nearest Interpolation

```
nearest_interp['Ozone'].plot(color='red', marker='o', linestyle='dotted', figsize=(30, 5))  
airquality['Ozone'].plot(title='Ozone', marker='o')
```



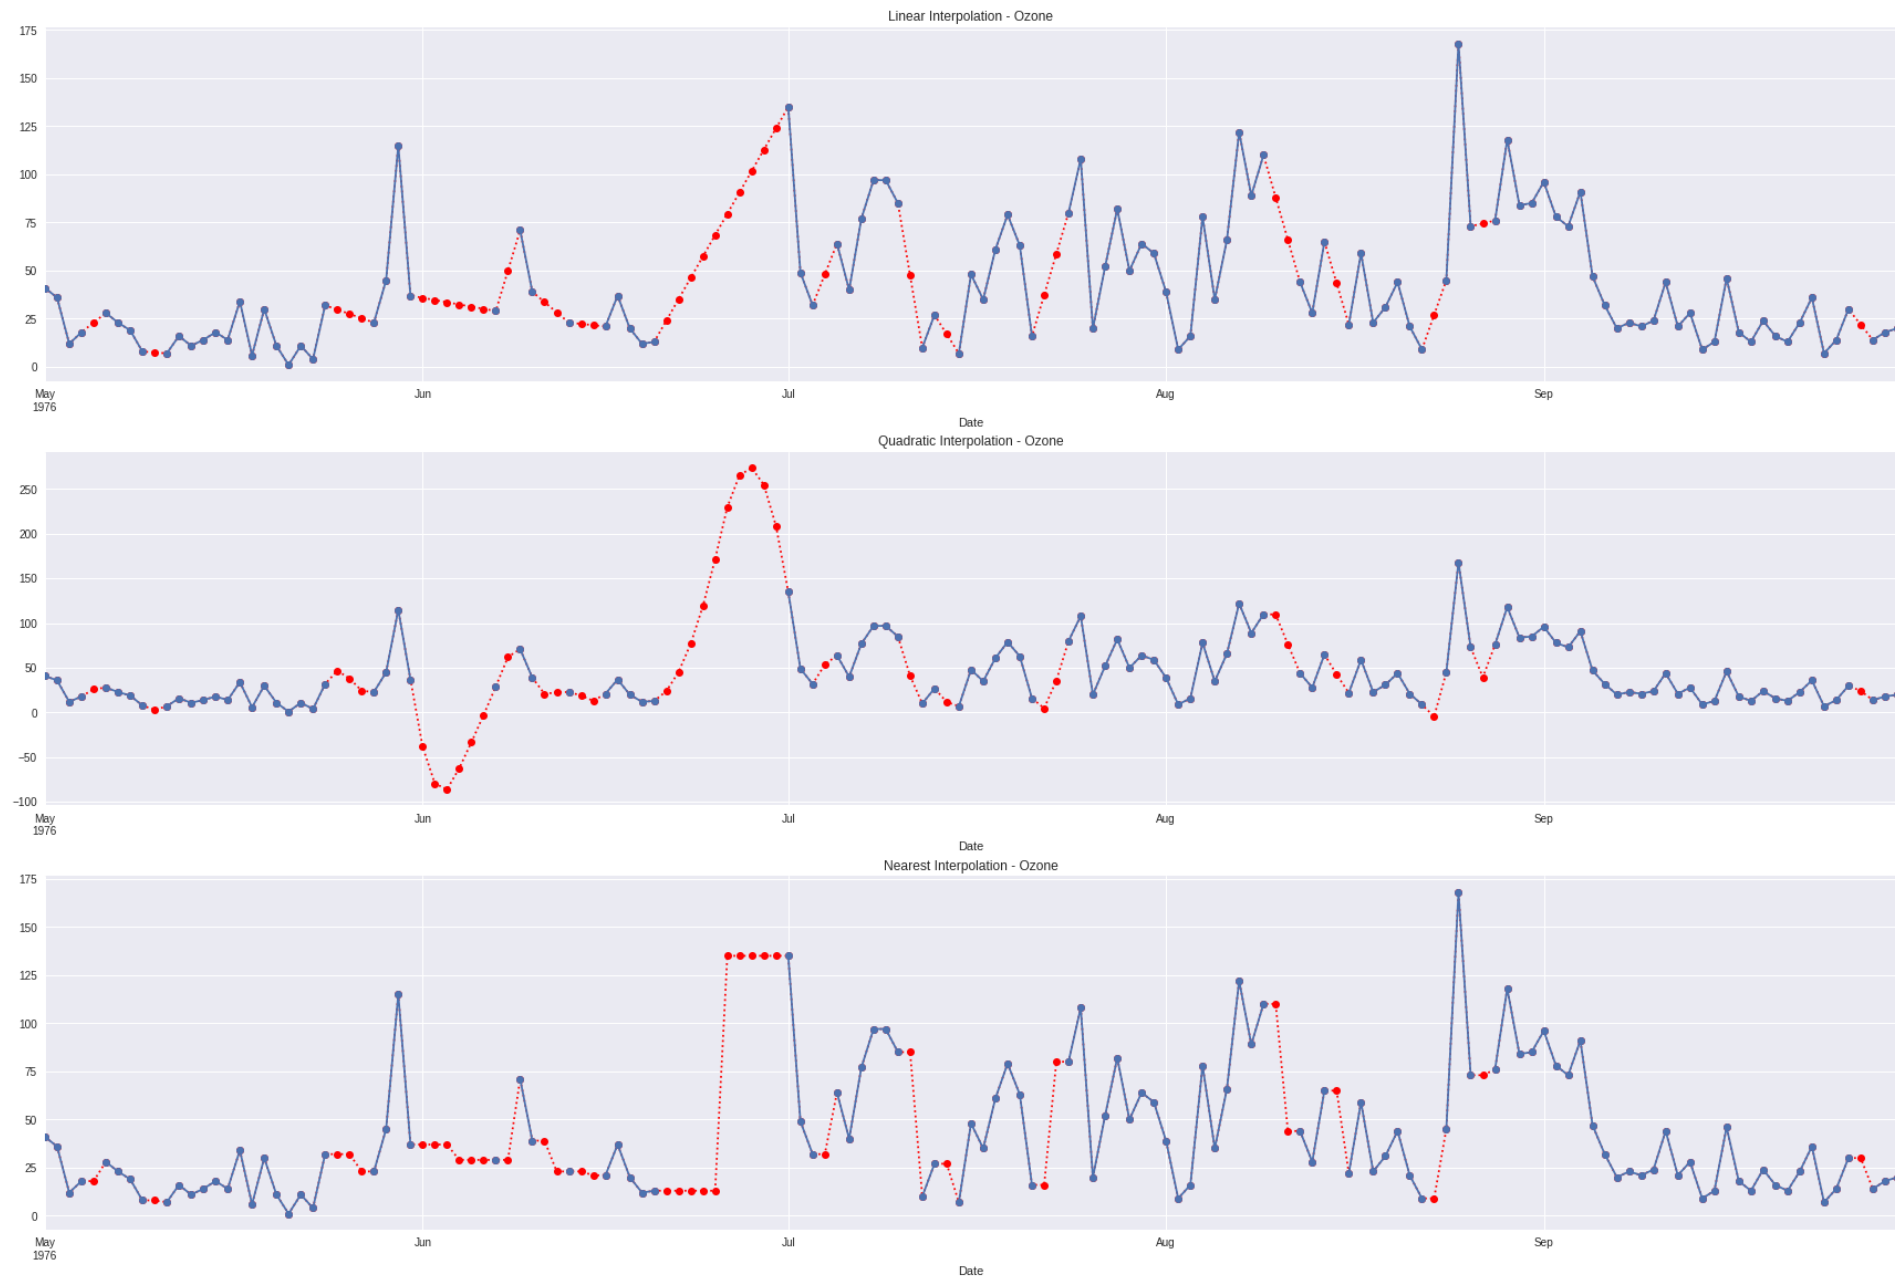
A comparison of the interpolations

```
# Create subplots
fig, axes = plt.subplots(3, 1, figsize=(30, 20))

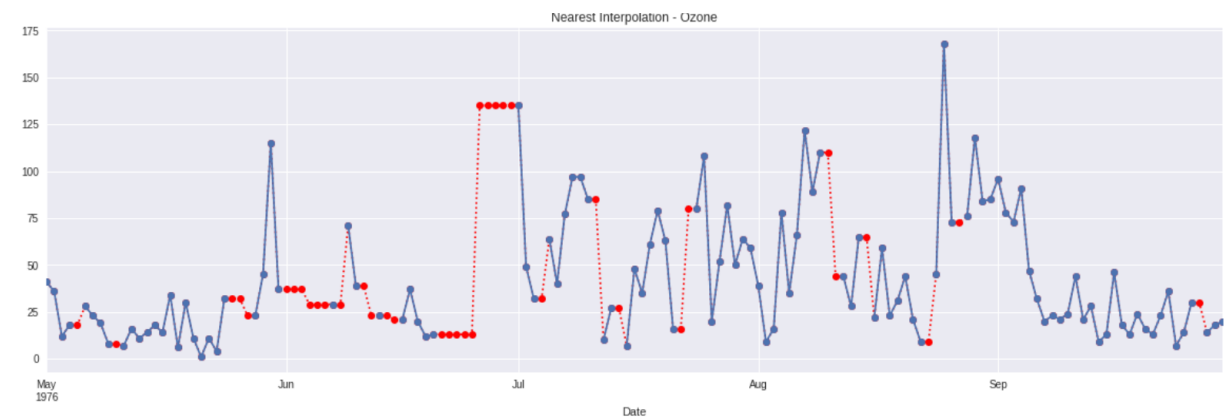
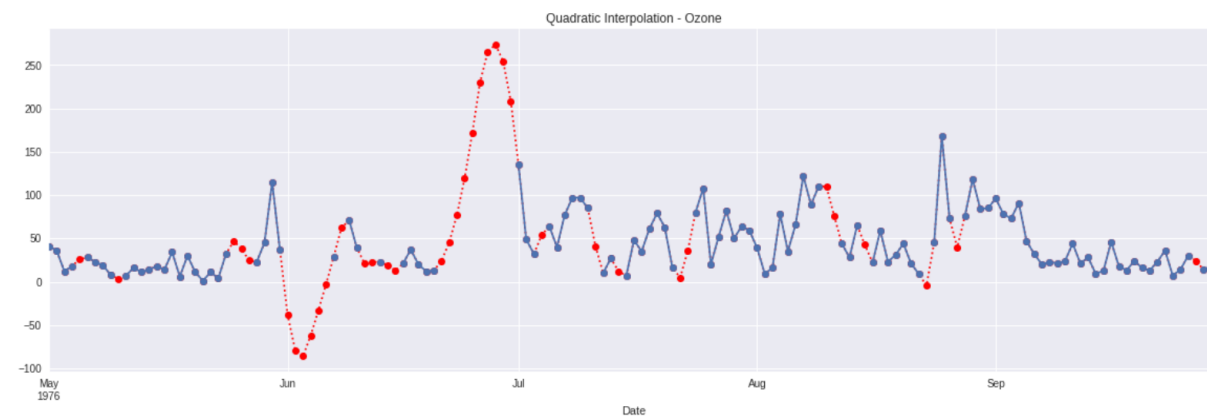
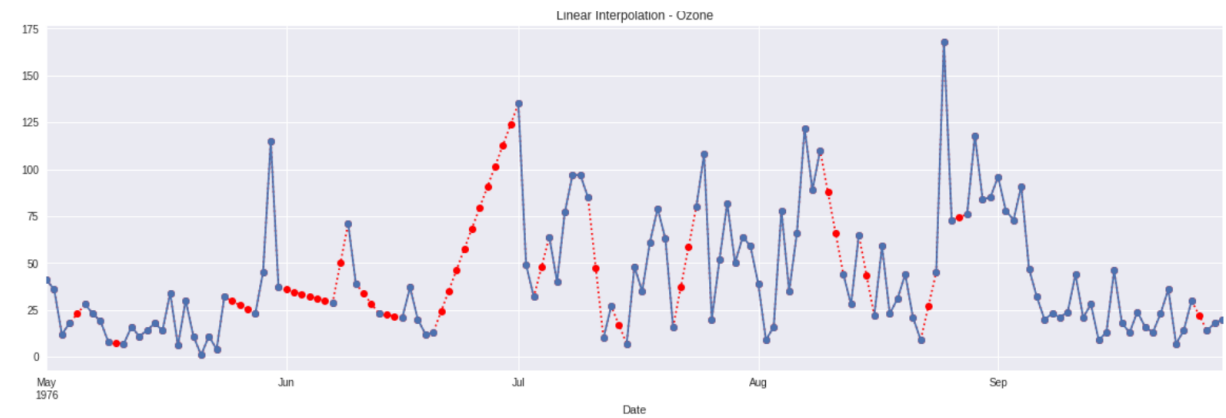
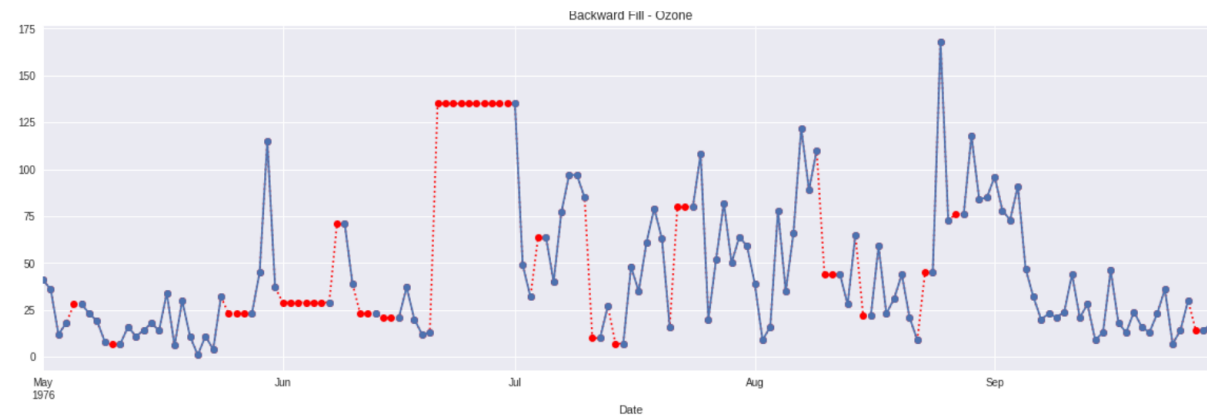
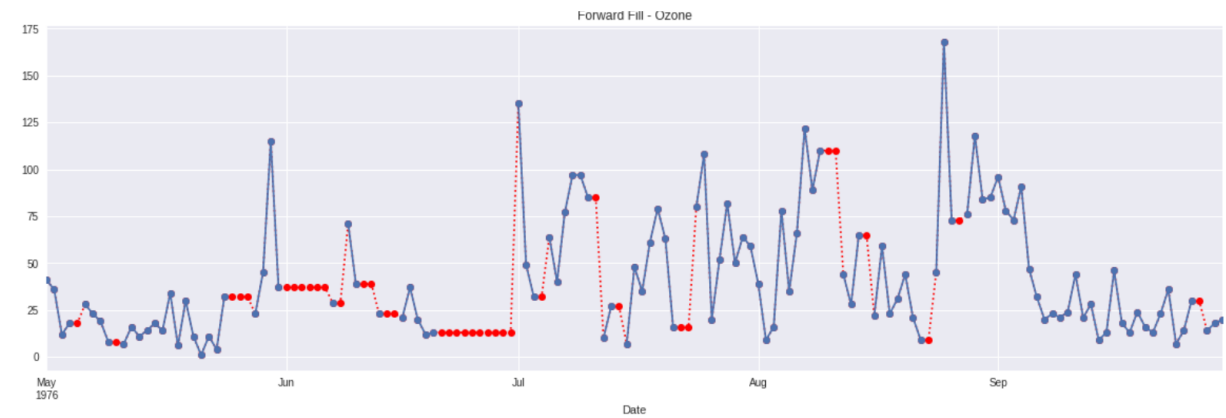
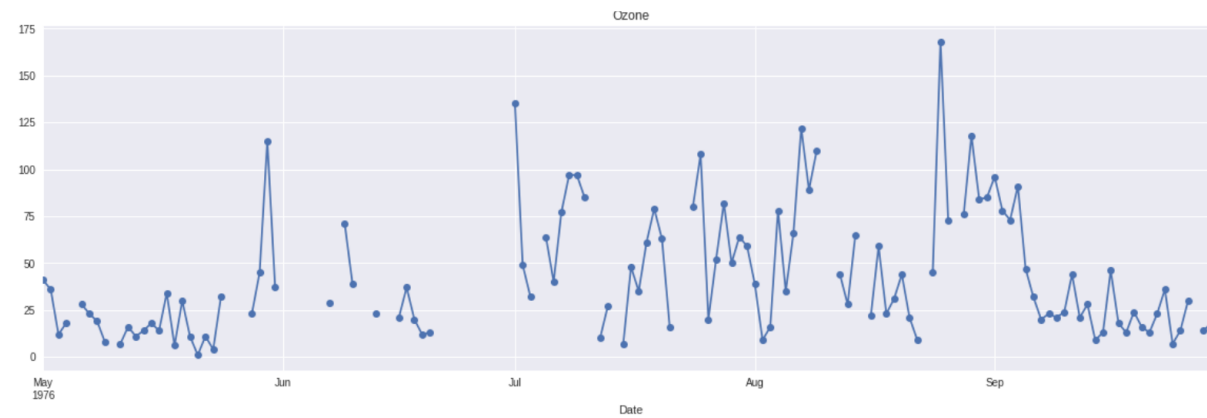
# Create interpolations dictionary
interpolations = {'Linear Interpolation': linear_interp,
                  'Quadratic Interpolation': quadratic_interp,
                  'Nearest Interpolation': nearest_interp}

# Visualize each interpolation
for ax, df_key in zip(axes, interpolations):
    interpolations[df_key].Ozone.plot(color='red', marker='o',
                                     linestyle='dotted', ax=ax)
    airquality.Ozone.plot(title=df_key + ' - Ozone', marker='o', ax=ax)
```

A comparison of the interpolations



A comparison of imputation techniques



Summary

- Time-series plot of imputed DataFrame
- Comparison of imputations

Let's practice!

DEALING WITH MISSING DATA IN PYTHON